

```
In [61]: !jupyter nbconvert --to html /content/CS401Lab3b.ipynb
```

```
[NbConvertApp] Converting notebook /content/CS401Lab3a.ipynb to html  
[NbConvertApp] Writing 1113260 bytes to /content/CS401Lab3a.html
```

```
In [56]: from google.colab import drive  
drive.mount('/content/drive')
```

#https://www.kaggle.com/datasets/ruiromanini/mtcars/ (where I got the dataset)

Drive already mounted at /content/drive; to attempt to forcibly remount, call `drive.mount("/content/drive", force_remount=True)`.

1. Data Exploration (10 points):

(a) Load the "mtcars" dataset and describe its structure, including the number of observations and variables.

(b) Explore the dataset by calculating summary statistics and visualizing the data. Create scatter plots to examine the relationships between the independent variables and the target variable (mpg).

```
In [55]: #a) Load the "mtcars" dataset and describe its structure, including the number of c  
import pandas as pd  
from sklearn import datasets  
  
file_path = '/content/drive/MyDrive/mtcars.csv'  
df = pd.read_csv(file_path)  
  
print(df)  
print(df.shape[0])  
print(df.shape[1])
```

	model	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	\
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	
5	Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	
6	Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	
7	Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	
8	Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	
9	Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	
10	Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	
11	Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	
12	Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	
13	Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	
14	Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	
15	Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	
16	Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	
17	Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	
18	Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	
19	Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	
20	Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	
21	Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	
22	AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	
23	Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	
24	Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	
25	Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	
26	Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	
27	Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	
28	Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	
29	Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	
30	Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	
31	Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	

	gear	carb
0	4	4
1	4	4
2	4	1
3	3	1
4	3	2
5	3	1
6	3	4
7	4	2
8	4	2
9	4	4
10	4	4
11	3	3
12	3	3
13	3	3
14	3	4
15	3	4
16	3	4
17	4	1
18	4	2
19	4	1
20	3	1
21	3	2
22	3	2
23	3	4
24	3	2
25	4	1
26	5	2
27	5	2
28	5	4

```

29     5     6
30     5     8
31     4     2
32
12

```

a) Load the "mtcars" dataset and describe its structure, including the number of observations and variables.

There are 32 rows in total with 11 columns made up of: model, mpg, cyl, disp, hp, drat, wt, qsec, vs, am, gear, carb.

This is the model of the car, miles per gallon, number of cylinders in the engine, displacement (Engine's volume), horsepower, drat (affects car's acceleration and top speed), weight of the car, qsec is the amount of time it takes to go from 0mph to 60mph, vs is the engine type, am is the transmission type, gear is gears & carb is the number of carburetors that the engine has.

```

In [57]: #b) Explore the dataset by calculating summary statistics and visualizing the data.
#Create scatter plots to examine the relationships between the independent variables
import seaborn as sns
import matplotlib.pyplot as plt

summary_stats = df.describe()
print(summary_stats)

sns.heatmap(df.corr(), annot=True, cmap='coolwarm', square=True)
plt.show()

for column in df.columns:
    if column != 'mpg': # Exclude 'mpg' from independent variables
        plt.figure(figsize=(10, 10))
        sns.scatterplot(x=column, y='mpg', data=df)
        plt.title(f'Scatter Plot of {column} vs. mpg')
        plt.xlabel(column)
        plt.ylabel('mpg')
        plt.show()


```


	mpg	cyl	disp	hp	drat	wt \
count	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000
mean	20.090625	6.187500	230.721875	146.687500	3.596563	3.217250
std	6.026948	1.785922	123.938694	68.562868	0.534679	0.978457
min	10.400000	4.000000	71.100000	52.000000	2.760000	1.513000
25%	15.425000	4.000000	120.825000	96.500000	3.080000	2.581250
50%	19.200000	6.000000	196.300000	123.000000	3.695000	3.325000
75%	22.800000	8.000000	326.000000	180.000000	3.920000	3.610000
max	33.900000	8.000000	472.000000	335.000000	4.930000	5.424000


	qsec	vs	am	gear	carb
count	32.000000	32.000000	32.000000	32.000000	32.0000
mean	17.848750	0.437500	0.406250	3.687500	2.8125
std	1.786943	0.504016	0.498991	0.737804	1.6152
min	14.500000	0.000000	0.000000	3.000000	1.0000
25%	16.892500	0.000000	0.000000	3.000000	2.0000
50%	17.710000	0.000000	0.000000	4.000000	2.0000
75%	18.900000	1.000000	1.000000	4.000000	4.0000
max	22.900000	1.000000	1.000000	5.000000	8.0000


```
<ipython-input-57-53a29aa8611e>:9: FutureWarning: The default value of numeric_only
in DataFrame.corr is deprecated. In a future version, it will default to False.
Select only valid columns or specify the value of numeric_only to silence this warning.
```


```
sns.heatmap(df.corr(), annot=True, cmap='coolwarm', square=True)
```
































b) Explore the dataset by calculating summary statistics and visualizing the data. Create scatter plots to examine the relationships between the independent variables and the target variable (mpg).

The summary statistics tell us information about each of the 12 columns. This includes, the count, mean, std, min, max & number at each quartile.

We also have a heatmap which shows which variables have the most correlation. The darker red it is, the higher the correlation. The darker the blue, the lower amount of correlation.

This is important to do as it allows us to see the relationship between the other independent variables and MPG before the scatterplots are made. Drat has the highest amount of correlation with mpg and cyl has the least amount of correlation of

Positive/Negative Linear Relationship No Linear Relationship Outliers Clusters or Groups
Correlation Strength Non-Linear Relationships D

Cyl, vs, am, gear & cabs are split up into groups of whole numbers. Carb is mainly split up into 1, 2, 3 & 4 but there is an outlier in both 6 & 8. There's a strong negative linear relationship in (cyl vs. mpg), (disp vs. mpg), (hp vs. mpg), (wt vs. mpg) & (carb vs. mpg)

(qsec vs. mpg) & (drat vs. mpg) are the only ones with a notable with strong positive correlation while (am vs. mpg) & (vs vs. mpg) does have some correlation.

(model vs. mpg) has no apparent linear relationship.

1. Simple Linear Regression (30 points):

- Select one independent variable from the "mtcars" dataset that you believe may have a strong linear relationship with the target variable (mpg).
- Implement a simple linear regression model to predict mpg using the selected independent variable.
- Calculate the model's coefficients (slope and intercept) and evaluate its performance using appropriate regression evaluation metrics (on testing dataset).

In []: *#a) Select one independent variable from the "mtcars" dataset that you believe may have a strong linear relationship with mpg so I will choose qsec vs. mpg*

In []: *#b) Implement a simple linear regression model to predict mpg using the selected independent variable
#c) Calculate the model's coefficients (slope and intercept) and evaluate its performance*

```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score

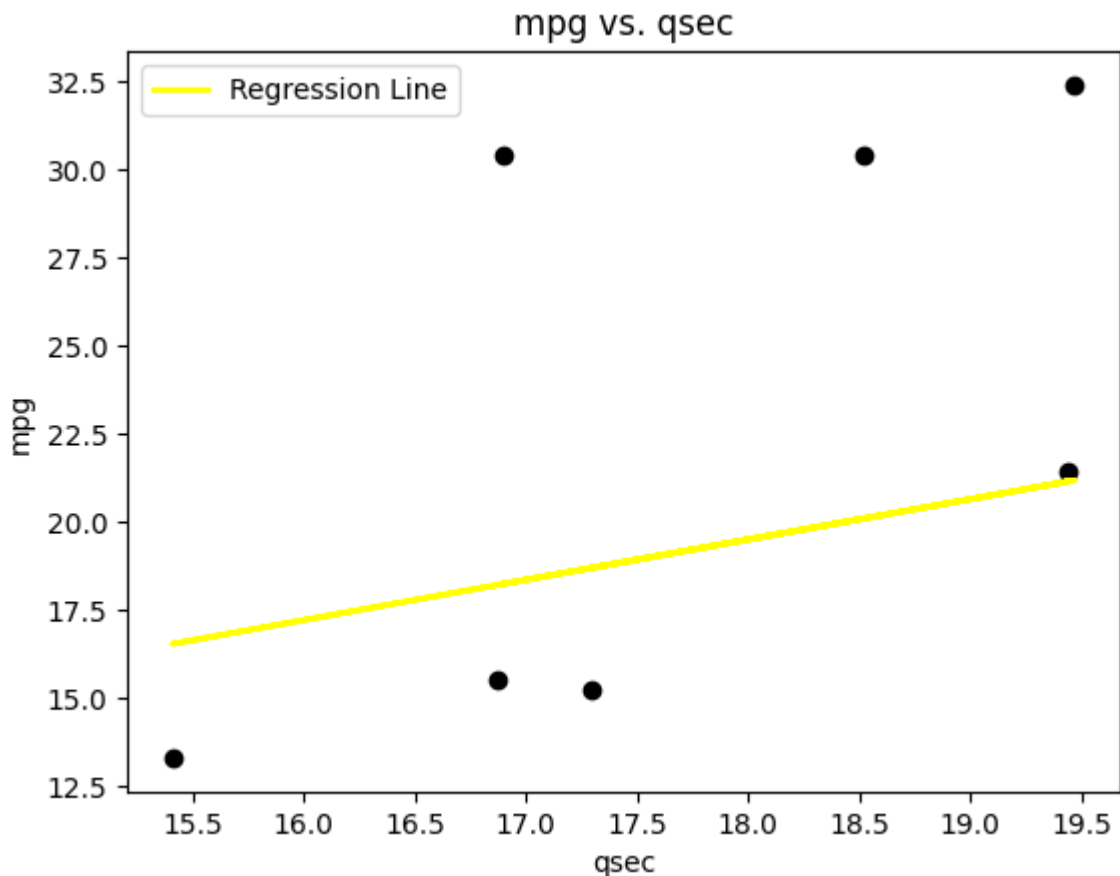
x = df[['qsec']]
y = df['mpg']
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)

model = LinearRegression().fit(x_train, y_train)
Prediction = model.predict(x_test)
slope = model.coef_[0]
intercept = model.intercept_
mse = mean_squared_error(y_test, Prediction)
r2 = r2_score(y_test, Prediction)

print(f"Slope: {slope:.2f}")
print(f"Intercept: {intercept:.2f}")
print(f"Mean Squared Sum of Errors: {mse:.2f}")
print(f"R2: {r2:.2f}")

plt.scatter(x_test, y_test, color='Black')
plt.plot(x_test, Prediction, color='Yellow', linewidth=2, label='Regression Line')
plt.xlabel('qsec')
plt.ylabel('mpg')
plt.legend()
plt.title('mpg vs. qsec')
plt.show()
```

Slope: 1.15
Intercept: -1.12
Mean Squared Sum of Errors: 58.58
R2: 0.00



b) Implement a simple linear regression model to predict mpg using the selected independent variable. c) Calculate the model's coefficients (slope and intercept) and evaluate its performance using appropriate regression evaluation metrics (on testing dataset).

This is done! We can see that the slope is 1.15, the intercept is -1.12 and the MSE is 58.58.

This is similar to what was asked of us last week to it made doing this question a lot easier.

1. Multiple Linear Regression (40 points):

(a) Implement a multiple linear regression model using a combination of independent variables from the "mtcars" dataset.

(b) Train the model to predict mpg using multiple features.

(c) Evaluate the model's performance using appropriate regression evaluation metrics (on testing dataset).

```
In [43]: #(a) Implement a multiple linear regression model using a combination of independent
#(b) Train the model to predict mpg using multiple features.
#On a side note, it's when we are doing stuff like this that I just wish that we we
#a repeat of what was done previously.

X = df[['qsec', 'wt', 'hp']]
Y = df['mpg']
model = LinearRegression()
model.fit(X,Y)

print(f"Slope: {model.coef_}")
print(f"Intercept: {model.intercept_}")
```

Slope: [0.51083369 -4.3587972 -0.01782227]
 Intercept: 27.610526858205063

Something of note is `X = df[['qsec', 'wt', 'hp']]` is made up of the dataframe and array. The two `[]` need to be used here or it simply doesn't work.

Apart from that, it's simple to do the Multiple Linear Regression model.

```
In [52]: #(c) Evaluate the model's performance using appropriate regression evaluation metrics
from math import sqrt

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = LinearRegression()
model.fit(X_train, y_train)
PredictionB = model.predict(X_test)
RMSE = sqrt(mean_squared_error(y_test, PredictionB))


print(f"Mean Squared Error (MSE): {mean_squared_error(y_test, PredictionB):.2f}")
print(f"R-squared (R2): {r2_score(y_test, PredictionB):.2f}")
print(f"Root Mean Squared Error (RMSE): {RMSE:.2f}")

plt.scatter(y_test, PredictionB)
plt.xlabel("Actual Values")
plt.ylabel("Predicted Values")
#plt.legend()
plt.show()
```

Mean Squared Error (MSE): 12.89

R-squared (R2): 0.78

Root Mean Squared Error (RMSE): 3.59



Mean Squared Error (MSE): 12.89 R-squared (R2): 0.78 Root Mean Squared Error (RMSE): 3.59

We get the expected result as seen in the scatterplot. The correlation is positive. There is a small cluster at the top right of the model that should be made note of but our actual results are not unexpected.

1. Discussion and Conclusion 20 points):

(a) Compare the performance and interpretability of the simple linear regression model with the multiple linear regression model. Discuss the trade-offs between simplicity and complexity.

The simple linear regression model shows the effect 2 variables have on each other. It shows us the relationship between the 2 variables. it's easy to produce and easy to interpret.

Compared to Simple Linear Regression, Multiple Linear Regression shows the complexities of the linear relationships by comparing to more variables. While this can lead to a better model performance, too many variables could make the model moot.

Simply, the Simple Linear Regression model is a lot easier to understand. We are comparing 2 variables to each other. Multilinear Regression Model has more factors to influence the model, thus makes it more complicated.

For the most accurate answer, the multi-linear regression model is better for predicting how the relationship is but if there are variables that do not matter, it can skew the model.

Choosing between the two depends on the situation. Knowing the best one to choose depending on the situation is key to a successful model.

(b) Reflect on the insights gained from the assignment and the implications for predicting fuel efficiency in car models.

Simple & Multiple Linear Regression models can be used for predicting fuel efficiency in car models in the future. However, the variables being compared must be important. By developing regression models to predict fuel efficiency, car makers/manufacturers can make a more informed decision when making the car. Seeing how different fuel types react affect sales of cars could be very important. E.G: Electricity vs. Gas.

This assignments shows the importance of fuel efficiency, linear regression models and why you should use both simple & multiple linear regression models.